

Research and Development of an Automated System for Concept Drift Detection and Adaptive Machine Learning Model Updating

Master's Thesis

Student: Le Phuc Duc

Supervisor 1: Assoc.Prof.Dr. Thoai Nam

Supervisor 2: Dr. Nguyen Quang Hung

Ho Chi Minh City University of Technology
Faculty of Computer Science and Engineering

2026



Presentation Outline

- 1 Introduction
- 2 Background: The MMD Drift Signal
- 3 Layer 1: Detection
- 4 Layer 2: Classification
- 5 Layer 3: Adaptation
- 6 Evaluation & Results
- 7 System Deployment
- 8 Conclusion

Agenda

- 1 Introduction
- 2 Background: The MMD Drift Signal
- 3 Layer 1: Detection
- 4 Layer 2: Classification
- 5 Layer 3: Adaptation
- 6 Evaluation & Results
- 7 System Deployment
- 8 Conclusion







The Core Problem: Concept Drift

Mathematical Definition

$$P_{\text{train}}(X, Y) \neq P_{\text{online},t}(X, Y)$$
$$P(X, Y) = P(X)P(Y|X)$$

Real-world Impact:

- Banking Fraud (False Alarms $\uparrow$)
- Predictive Maintenance (Accuracy $\downarrow$)







Concept Drift: Distribution vs Pattern

Virtual Drift:

Only $P(X)$ changes.

Real Drift:

Decision boundary $P(Y|X)$ changes.




Objectives & System Architecture

A three-layer pipeline:


- **L1 — Detect** (*When?*)
ShapeDD-IDW: fast, unsupervised, $\alpha = 0.05$ calibrated.
- **L2 — Classify** (*What?*)
SE-CDT: label-free shape analyzer, 5 drift types.
- **L3 — Adapt** (*How?*)
Type-aware router: reset / warm-start / reuse.

Goals

Fast, reliable detection; unsupervised drift typing (no labels); end-to-end validation.



Related Works: The Gap in Current Literature



Evaluated Datasets: SEA, RTG, RBF, Sine [1, 2, 3]

Agenda

- 1 Introduction
- 2 Background: The MMD Drift Signal
- 3 Layer 1: Detection
- 4 Layer 2: Classification
- 5 Layer 3: Adaptation
- 6 Evaluation & Results
- 7 System Deployment
- 8 Conclusion



MMD: The Unsupervised Drift Signal

Why unsupervised?

- Supervised detection needs ground-truth labels right after each prediction.
- In real application, labels are delayed or expensive → monitor the **raw distribution** $P(X)$ instead.


Maximum Mean Discrepancy: map points into an RKHS and compare their mean embeddings (“centers of mass”).

MMD Statistic

$$\widehat{\text{MMD}}^2 = \underbrace{\text{Within } P}_{\text{self-sim}} + \underbrace{\text{Within } Q}_{\text{self-sim}} - \underbrace{\text{Cross } P \leftrightarrow Q}_{\text{dis-sim}}$$

$\widehat{\text{MMD}}^2 \approx 0$: same distribution; $\widehat{\text{MMD}}^2 \gg 0$: **drift candidate**.

MMD Detects Variance Shifts (Unlike Mean-based Tests)



From MMD Trace to ShapeDD's Triangle

Stream $\rightarrow$ trace:


- A sliding window of size $2h_1$ splits into W_{ref} and W_{test} ; their MMD at each step forms a continuous trace $\sigma(t)$.
- Stable data $\rightarrow$ small noise floor; crossing a drift $\rightarrow$ $\sigma(t)$ rises into a peak.

ShapeDD's insight — the triangle:

- A sudden drift yields a **triangular** $\sigma(t)$ (linear contamination as the window crosses the drift) — detect that *shape*, not a raw threshold.
- Peak height $\sigma(t_c) =$ **drift magnitude** $\approx \text{MMD}^2(P, Q)$.

Bottleneck

Significance via a **permutation test** rebuilds the kernel thousands of times. We replace it with **IDW weighting** + **Gamma-null** (next).



Agenda

- 1 Introduction
- 2 Background: The MMD Drift Signal
- 3 Layer 1: Detection**
- 4 Layer 2: Classification
- 5 Layer 3: Adaptation
- 6 Evaluation & Results
- 7 System Deployment
- 8 Conclusion



Inverse Density-Weighted MMD (IDW-MMD)

Uniform weights ($1/n^2$):

- Dense majority clusters dominate the statistic.
- Early drifts at distribution **boundaries** get washed out.

IDW weighting (the idea)

Weight each point by **inverse local density**:


$$w_i \propto \frac{1}{\sqrt{d_i + \epsilon}}, \quad d_i = \sum_j k(x_i, x_j)$$

Dense core $\rightarrow$ small w_i ; sparse **boundary** $\rightarrow$ large w_i .

Full IDW-MMD Statistic

$$\widehat{\text{MMD}}_{\text{IDW}}^2 = \underbrace{\sum_{i \neq j} W_{ij}^{XX} k(x_i, x_j)}_{\text{Within } P \text{ (Weighted)}} + \underbrace{\sum_{p \neq q} W_{pq}^{YY} k(y_p, y_q)}_{\text{Within } Q \text{ (Weighted)}} - 2 \underbrace{\sum_{i,p} \frac{1}{nm} k(x_i, y_p)}_{\text{Cross } P \leftrightarrow Q \text{ (Uniform)}}$$

- **Inverse Density** (W_{ij}^{XX}, W_{pq}^{YY}): amplifies boundary points.
- **Uniform weight** (XY term): anchors geometric overlap.



Result: sensitive to early boundary changes, no extra false alarms on the bulk.


Statistical Calibration (Gamma-Null Approximation)

The Permutation Bottleneck:


- Standard ShapeDD runs $B = 2500$ permutations to estimate p-value.
- Each permutation re-builds the kernel matrix $\rightarrow$ extremely slow.

Gamma-Null Approximation Solution:

- Under H_0 , the MMD statistic is approximately Gamma-distributed.
- We estimate the first two moments (mean & variance) from a small number of bootstrap shuffles ($B = 20$), reusing the kernel matrix.
- Matches accuracy while achieving a **7–10 $\times$** end-to-end speedup.



ShapeDD-IDW in Action: MMD Trace $\rightarrow$ Single Drift Event



The detector, end-to-end:

- 1 (a) Stream $\rightarrow$ compute the IDW-MMD trace $\sigma(t)$.
- 2 (b) Isolate peaks: smooth $\sigma(t)$, keep points above $\bar{\sigma} + \text{std}$. Ignore blips under threshold.
- 3 (c) Validate & group: the Gamma-null p -value (Bonferroni-adjusted) confirms the peak; the whole super-threshold burst collapses to one event within tolerance δ .

Outcome

One calibrated detection per true drift — no permutation test, few false alarms.

Agenda

- 1 Introduction
- 2 Background: The MMD Drift Signal
- 3 Layer 1: Detection
- 4 Layer 2: Classification**
- 5 Layer 3: Adaptation
- 6 Evaluation & Results
- 7 System Deployment
- 8 Conclusion



CDT-MSW — Supervised Classification Baseline

CDT-MSW (Guo et al., 2022) groups drift into:


- **TCD (Transient)**: Sudden, Blip, Recurrent.
- **PCD (Progressive)**: Gradual, Incremental.

Growth Process:

- Tracks the variance of accuracy degradation over a multi-sliding window.
- Short degradation → TCD; sustained degradation → PCD.

Limitation

CDT-MSW needs **ground-truth labels** at every step to compute the accuracy curve — impractical for data streams where labels arrive late or never.



Motivation for SE-CDT


Replace the accuracy curve with the *shape* of the MMD signal $\sigma(t)$. The same TCD/PCD distinction can then be made **without any labels at runtime**.

SE-CDT: Reading the MMD Shape (9 Attributes)


Type drift without labels by reading the shape of $\sigma(t)$ at the detection point:

- **Geometric:** n_p (peak count), WR (width ratio), CV (coeff. of variation), SNR, PPR, DPAR.
- **Temporal:** LTS (linear trend), SDS (step score), MS (monotonicity).
- + **VR:** one *raw-data* signal (not from $\sigma(t)$); the slow-drift tie-breaker.

1. Peak-Level Geometry & Temporal Trends




2. Multi-Peak Geometry & Periodicity



SE-CDT: Variance Ratio (VR) for Slow Drifts






Challenge: Gradual and Incremental share flat MMD humps.

Solution: Variance Ratio (VR). Gradual (mixture) $\Rightarrow$ variance **spike**. Incremental (single shift) $\Rightarrow$ **flat**.



SE-CDT: Decision Logic & Drift Signatures

Simplified decision tree:




Hybrid Pipeline: Standard MMD to Trace, IDW-MMD to Validate

Why use two MMD variants instead of one?

- **Standard MMD:** Fast, preserves the natural geometry of the trace (gradual rising trends, duration) needed for downstream shape classification.
- **IDW-MMD:** Highly sensitive to shift, rigorously controls Type-I error ($\alpha = 0.05$). Density computation is costly, so it's only used for validation.

Two-Stage Architecture

- **Stage 1 (Trace):** Run **Standard MMD** continuously to find candidate peaks.
- **Stage 2 (Validate):** Compute **IDW-MMD** (p -val) only at peaks. If confirmed, pass the Standard MMD trace shape to the SE-CDT Classifier.



Agenda

- 1 Introduction
- 2 Background: The MMD Drift Signal
- 3 Layer 1: Detection
- 4 Layer 2: Classification
- 5 Layer 3: Adaptation**
- 6 Evaluation & Results
- 7 System Deployment
- 8 Conclusion



Optimizing Update Cost

Detected Drift Type	Adaptation Policy
Sudden	Full Reset
Incremental	Continuous Update
Gradual	Weighted Window
Recurrent	Model Reuse (from Memory)
Blip	Ignore (Do Nothing)

Goal

Minimize catastrophic forgetting and compute overhead by avoiding a "one-size-fits-all" retraining policy.

Kafka Integration

Triggers downstream microservices to hot-swap or fine-tune models dynamically.



Agenda

- 1 Introduction
- 2 Background: The MMD Drift Signal
- 3 Layer 1: Detection
- 4 Layer 2: Classification
- 5 Layer 3: Adaptation
- 6 Evaluation & Results**
- 7 System Deployment
- 8 Conclusion



14 Evaluation Datasets (Grouped):

- **Sudden** $P(X)$: Gaussian Shift (Moderate), GCS (3 concepts), $4 \times$ Random Uniform (Sensitivity tests: Mild to Ultra Severe).
- **Blip** $P(X)$: RBF Blips (short-term recurring centroid shifts).
- **Real Drift** $P(Y|X)$ (**Control Group**): Standard SEA, Rotating Hyperplane, LED Abrupt (no feature drift; to verify negative controls).
- **Semi-Synthetic**: Electricity Semi-Synthetic (real features, 10 injected drift points).
- **Slow & Cycle**: Gaussian Gradual (linear transition), GRS (Recurrent $A \leftrightarrow B$), Gaussian Stationary (H_0 calibration).

Experimental Protocol

- **Runs**: 30 independent seeds per stream to ensure statistical stability.
- **Layer 1 (Detection)**: Tolerance $\delta = 75$, cooldown = 150. Measured via F1-score, EDR, and False Positives.
- **Layer 2 (Classification)**: Oracle Mode (1800 events total: $6 \text{ scenarios} \times 30 \text{ seeds} \times 10 \text{ events}$) to measure CAT and SUB accuracy.
- **Layer 3 (Adaptation)**: Prequential accuracy (test-then-train) on Stepping and Sudden streams.

Classification Setup in Oracle Mode

Streams (5-D Gaussian, 1800 events total):

Scenario	Composition
Mixed A	Sudden $\rightarrow$ Gradual $\rightarrow$ Recurrent
Mixed B	Blip $\rightarrow$ Incremental $\rightarrow$ Sudden
Repeated Sudden	10 $\times$ Sudden
Repeated Gradual	10 $\times$ Gradual
Repeated Incremental	10 $\times$ Incremental
Repeated Recurrent	10 $\times$ Recurrent

6 scenarios $\times$ 30 seeds $\times$ 10 events = **1800 events**. Each non-recurrent event shifts a different feature subset, so new concepts are genuinely distinct.

Oracle mode


The classifier receives the ground-truth drift *position* and is scored only on the *typing* step. This isolates classification skill from detection skill — the same way the detection protocol ($\delta = 75$, cooldown) isolates detection.

Two metrics

CAT — TCD vs PCD (the decision that drives adaptation).
SUB — exact 5-class subtype (micro / macro).



Result 1a: Drift Detection Performance



Statistical Significance (Nemenyi Test):

- **Meaning:** Post-hoc pairwise comparison of detectors across all 14 datasets.
- **Critical Difference (CD):** Threshold of 2.81 ($\alpha = 0.05$).
- **Interpretation:** Groups connected by a thick line (like DAWIDD, MMD, and SE-CDT) have **no statistically significant difference** in performance.
- **Result:** SE-CDT (rank 3.89) is statistically tied with state-of-the-art DAWIDD (rank 3.46).

Result 1b: Type-I Error on Stationary Data

Type-I Error at nominal $\alpha = 0.05$

Distribution	Standard MMD (perm)	IDW-MMD (Gamma null)	SE-CDT (composite)
Gauss-iid d=5	0.025 $\pm$ 0.022	0.040 $\pm$ 0.027	0.115 $\pm$ 0.044
Gauss-iid d=10	0.055 $\pm$ 0.032	0.050 $\pm$ 0.030	0.080 $\pm$ 0.038
Gauss-AR1 d=5	0.045 $\pm$ 0.029	0.075 $\pm$ 0.037	0.140 $\pm$ 0.048

Evaluation on Stationary Streams:

- Evaluated across 3 stationary datasets (no drift) over 30 runs to measure false alarms.
- **IDW-MMD Calibration:** Holds extremely close to the nominal rate ($\alpha = 0.05$), ranging between 0.025–0.075.
- **Composite Trace (SE-CDT):** Becomes anti-conservative (0.080–0.140) due to peak selection and sliding-window dependency.

Result 1c: Comprehensive Detection Performance

Overall Metrics across all evaluation streams:

Method	Precision	Recall (EDR)	F1-Score	Delay	False Pos.
DAWIDD	<u>0.435</u>	0.734	0.532	38	10.3
IDW_MMD	0.432	0.737	<u>0.531</u>	36	<u>9.6</u>
SE_CDT	0.432	0.737	<u>0.531</u>	36	<u>9.6</u>
ShapeDD_IDW	0.432	0.737	<u>0.531</u>	36	<u>9.6</u>
MMD	0.429	0.734	0.527	35	10.5
ShapeDD	0.377	<u>0.749</u>	0.491	<u>34</u>	13.2
D3	0.558	0.472	0.489	16	0.6
KS	0.224	0.775	0.335	<u>34</u>	27.2

- **Synthetic $P(X)$ drifts** (Gaussian, Random Uniform): strong F1 on moderate-to-severe shifts (0.69–0.74), degrading gracefully on mildest case (0.51), with $FP \leq 10$.
- **Recurrent/Blip** (Rbfbliips, Stagger Recurrent): reliably detected ($F1 \approx 0.74$ –0.76) with minimal FP.
- **Real-drift controls** (SEA, Hyperplane, LED): low F1 *by design* — no $P(X)$ shift is present, confirming no false alarms on pure $P(Y|X)$ drift.



Result 1c: Per-Dataset Detection F1 (1/2)

Part 1: Basic & Gradual/Mild Drift

Method	Electricity	Gaussian Gradual Synthetic	Gaussian Moderate	Random Mild
DAWIDD	0.250	0.472	<u>0.753</u>	0.594
IDW_MMD	0.383	<u>0.470</u>	0.736	0.511
SE_CDT	0.383	<u>0.470</u>	0.736	0.511
ShapeDD_IDW	0.383	<u>0.470</u>	0.736	0.511
MMD	<u>0.291</u>	0.462	0.740	<u>0.566</u>
ShapeDD	0.267	0.451	0.700	0.529
D3	0.213	0.029	0.998	0.012
KS	0.263	0.315	0.356	0.335

Part 2: Random Uniform & Hyperplane

Method	Random Moderate	Random Severe	Random Ultra Severe	Hyperplane
DAWIDD	0.728	<u>0.739</u>	<u>0.738</u>	<u>0.136</u>
IDW_MMD	<u>0.688</u>	0.708	0.712	0.179
SE_CDT	<u>0.688</u>	0.708	0.712	0.179
ShapeDD_IDW	<u>0.688</u>	0.708	0.712	0.179
MMD	0.685	0.707	0.709	0.169
ShapeDD	0.619	0.636	0.638	0.197
D3	0.244	0.931	0.933	0.000
KS	0.407	0.434	0.434	0.240



Result 1c: Per-Dataset Detection F1 (2/2)

Part 3: Led, Rbfbliips & Stagger


Method	Led Abrupt	Rbfbliips	Stagger	Stagger Recurrent Explicit
DAWIDD	0.130	<u>0.749</u>	0.750	0.749
IDW_MMD	0.148	0.744	<u>0.755</u>	<u>0.753</u>
SE_CDT	0.148	0.744	<u>0.755</u>	<u>0.753</u>
ShapeDD_IDW	0.148	0.744	<u>0.755</u>	<u>0.753</u>
MMD	0.127	0.747	0.742	0.739
ShapeDD	0.163	0.670	0.681	0.679
D3	0.000	1.000	0.998	1.000
KS	<u>0.071</u>	0.371	0.459	0.458

Part 4: Sea & Overall Mean F1

Method	Standard Sea	Mean
DAWIDD	0.129	0.532
IDW_MMD	<u>0.118</u>	<u>0.531</u>
SE_CDT	<u>0.118</u>	<u>0.531</u>
ShapeDD_IDW	<u>0.118</u>	<u>0.531</u>
MMD	0.163	0.527
ShapeDD	0.151	0.491
D3	0.000	0.489
KS	0.206	0.335



Result 1d: Runtime Efficiency (1/2)



Result 1d: Runtime Efficiency (2/2)


Execution Time per Stream

Method	Runtime (s/stream)	Throughput (samples/s)	Speedup vs ShapeDD
KS	0.29	33,990	14.47×
D3	0.31	32,300	13.75×
ShapeDD-IDW	0.59	16,858	7.18×
IDW-MMD (standalone)	0.59	16,849	7.17×
SE-CDT	0.60	16,764	7.14×
ShapeDD (original)	4.26	2,349	1.00×
MMD	4.31	2,322	0.99×
DAWIDD	5.66	1,767	0.75×

Computational Speedup

Gamma-null approximation makes SE-CDT roughly **7×** **faster** than permutation-based ShapeDD (**0.60 s** vs. **4.26 s** per 10k samples) and standard MMD (4.31 s), bringing kernel-based detection into the latency class of lightweight tests like KS and D3.

Result 2a: Unsupervised Classification (SE-CDT in Oracle Mode)



Two enhancements added

Concept Memory: ring buffer of post-drift snapshots (Standard-MMD match) → recovers Recurrent typing.
Online Self-Calibration: rolling-quantile baseline adapts SNR/WR/CV thresholds → robustness on noisy streams.

Headline — label-free typing

CAT = 80.1% (TCD vs PCD — the decision that drives adaptation), at **92.6% recall**, with **no labels**.

Per-type subtype accuracy (1800 events):


Drift Type	Acc.	Notes
Sudden	82.4%	Sharp, well-separated peaks
Recurrent	71.5%	Enabled by Concept Memory
Blip	60.8%	Multi-peak tolerance rule
Incremental	30.3%	Overlaps Gradual on VR
Gradual	27.2%	57% absorbed into Recurrent

Honest limitation


SUB = 55.3%: slow types overlap on Variance Ratio — the weak point. Supervised CDT-MSW scores higher SUB but needs labels.



Result 2b: Confusion Matrix



Result 3: Model Adaptation (Normal / Streaming Mode)




Strategy	Step.	Sudd.
Type-Specific (Green)	71.4%	70.8%
Simple Retrain (Blue)	63.0%	62.2%
Periodic Retrain (Orange)	64.0%	63.7%
No Adaptation (Red)	54.8%	53.7%


Why Type-Specific helps

- **Avoids Waste:** Resetting on gradual drift discards history; periodic retrains ignore signals.
- **Optimal Update:** Sudden → full reset, Gradual → weighted-window (warm-start).
- **Gains: +16.6 pp (Stepping) and +17.1 pp (Sudden) vs No Adaptation ($p < 10^{-15}$).**

Result 3b: Adaptation on Sudden Drift



Result 3c: Adaptation on Mixed Drift Scenario




Agenda

- 1 Introduction
- 2 Background: The MMD Drift Signal
- 3 Layer 1: Detection
- 4 Layer 2: Classification
- 5 Layer 3: Adaptation
- 6 Evaluation & Results
- 7 System Deployment**
- 8 Conclusion



End-to-End Deployment Prototype on Apache Kafka



Drift at sample 1500 → accuracy dips → SE-CDT fires ($p \approx 10^{-6}$); CPU/memory stay flat (backup).

Agenda

- 1 Introduction
- 2 Background: The MMD Drift Signal
- 3 Layer 1: Detection
- 4 Layer 2: Classification
- 5 Layer 3: Adaptation
- 6 Evaluation & Results
- 7 System Deployment
- 8 Conclusion



Contributions of This Thesis

- ① **Detection (Layer 1):** A calibrated, label-free detector — ShapeDD with IDW optimal weighting and a **Gamma-null** p-value that replaces the permutation test for speedup.
- ② **Classification (Layer 2): SE-CDT** — a fully *unsupervised* drift-type classifier that reads 9 geometric/temporal shape attributes (plus a Variance Ratio) of the MMD trace.
- ③ **Adaptation (Layer 3):** A **type-aware** routing policy mapping each drift type to its optimal update, validated end-to-end on an Apache Kafka prototype.



Summary

- ➊ **Detection:** ShapeDD-IDW with Gamma-null p-value provides a calibrated, label-free test with an end-to-end speedup.
- ➋ **Classification:** SE-CDT operates label-free, achieving high Category Accuracy (80.1%) and event recall (92.6%).
- ➌ **Adaptation:** Type-Specific routing yields substantial gains (+17.1 pp on Sudden) over no adaptation.
- ➍ **System:** Validated through an end-to-end Kafka prototype deployment.

Key Limitations

- **Subtype Confusion:** Gradual and Incremental shifts remain challenging (27.2% / 30.3% accuracy) due to overlap on the Variance Ratio.
- **Empirical Thresholds:** The Concept Memory threshold ($\mu = 0.15$) was chosen empirically based on the H0 noise floor.
- **Adaptation Evaluation Scope:** Tested heavily on sudden-dominated streams; robustness on purely gradual/incremental scenarios requires further study.

References I



M. Harries, "Splice-2 comparative evaluation: Electricity pricing," 1999.



W. N. Street and Y. Kim, "A streaming ensemble algorithm (sea) for large-scale classification," in *Proceedings of the seventh ACM SIGKDD international conference on Knowledge discovery and data mining*, 2001, pp. 377–382.



G. Hulten, L. Spencer, and P. Domingos, "Mining time-changing data streams," in *Proceedings of the seventh ACM SIGKDD international conference on Knowledge discovery and data mining*, 2001, pp. 97–106.



J. Gama, P. Medas, G. Castillo, and P. Rodrigues, "Learning with drift detection," in *Advances in Artificial Intelligence – SBIA 2004*, ser. Lecture Notes in Computer Science, vol. 3171. Springer Berlin Heidelberg, 2004, pp. 286–295.



M. Baena-García, J. del Campo-Ávila, R. Fidalgo-Merino, A. Bifet, R. Gavaldà, and R. Morales-Bueno, "Early drift detection method," in *Proceedings of the 4th ECML PKDD International Workshop on Knowledge Discovery from Data Streams (IWKDDs 2006)*, Berlin, Germany, 2006, pp. 77–86.



H. Guo, H. Li, Q. Ren, and W. Wang, "Concept drift type identification based on multi-sliding windows," *Information Sciences*, vol. 585, pp. 1–23, 2022.



Z. Wang and W. Wang, "Concept drift detection based on kolmogorov–smirnov test," in *Artificial Intelligence in China*. Springer Singapore, 2020, pp. 273–280.





F. Hinder, J. Brinkrolf, V. Vaquet, and B. Hammer, "A shape-based method for concept drift detection and signal denoising," in *2021 IEEE Symposium Series on Computational Intelligence (SSCI)*. IEEE, 2021, introduces ShapeDD algorithm for shape-based drift detection.

THANK YOU FOR LISTENING!

Q & A

Student: Le Phuc Duc

Supervisor 1: Assoc.Prof.Dr. Thoai Nam

Supervisor 2: Dr. Nguyen Quang Hung

Backup: SE-CDT Classification Algorithm

Algorithm: SE-CDT — Unsupervised Drift Type Classification

Require: Drift magnitude signal $\sigma(t)$ from Detection Module

Require: Raw data window W for Variance Ratio

Ensure: Drift type: {TCD-Blip, TCD-Sudden, TCD-Recurrent, PCD-Gradual, PCD-Incremental}

```
1:  $\sigma_s \leftarrow \text{GaussianFilter}(\sigma, \sigma_g)$ 
2:  $P \leftarrow \text{FindPeaks}(\sigma_s, \bar{\sigma}_s + \alpha \sigma_{std})$ 
3: Compute features:  $n_p \leftarrow |P|$ , WR, CV, SNR, PPR, DPAR, LTS, SDS, MS
4:  $VR \leftarrow \text{VarianceRatio}(W)$ 
5:  $\ell \leftarrow \text{EstimateDriftDuration}(\sigma_s)$ 
6: if  $n_p \leq \tau_n^{\text{blip}}$  and  $WR < \tau_{WR}^{\text{blip}}$  and  $\text{BlipProfile}(\dots)$  then
7:   return TCD-Blip
8: end if
9: if  $\ell > \tau_\ell$  then
10:   if  $\text{IsIncremental}(VR, LTS, \dots)$  then
11:     return PCD-Incremental
12:   else
13:     return PCD-Gradual
14:   end if
15: end if
16: if  $n_p \leq \tau_n^{\text{sud}}$  and  $WR < \tau_{WR}$  and  $\text{SNR} > \tau_{\text{SNR}}$  then
17:   return TCD-Sudden
18: end if
19: if  $n_p \geq \tau_n^{\text{rec}}$  and  $CV < \tau_{CV}$  and  $LTS < \tau_{LTS}$  then
20:   return TCD-Recurrent
21: end if
22: if  $\text{IsIncremental}(VR, LTS, \dots)$  then
23:   return PCD-Incremental
24: else
25:   return PCD-Gradual
26: end if
```

Backup: What the 9 Shape Attributes Mean

All measured on the smoothed MMD signal $\sigma_s(t)$; a peak = a point above $\bar{\sigma}_s + 0.3\sigma_{\text{std}}$.

Geometric — peak structure (“what it looks like”)

- n_p (peak count): 1 → Sudden; 2–3 → Blip; many → Recurrent.
- **WR** = FWHM/($2h_1$): small → sharp peak (Sudden); large → broad hump (PCD).
- **CV** = std/mean of inter-peak gaps: low → evenly spaced (Recurrent).
- **SNR** = max / median: high → one clean peak above the noise (Sudden).
- **PPR**: gap between 2 tallest peaks / length; small → two close peaks (Blip).
- **DPAR** = min / max of the 2 peak heights; near 1 → twin peaks (Blip).

Temporal — time behaviour (“how it evolves”)


- **LTS**: R^2 of a linear fit; high → straight ramp (Incremental).
- **SDS**: fraction of jumps > 1.5 std; high → step-like (Incremental).
- **MS** = $|\uparrow - \downarrow| / (\uparrow + \downarrow)$: high → one-directional (Incremental); low → oscillating (Gradual).

Quick guide

Sudden vs slow → **WR, SNR**; Blip → **PPR+DPAR**; Recurrent → n_p +**CV**; Incr. vs Grad. → **LTS/SDS/MS** (then **VR**).

The 6 geometric attributes answer *what the peak looks like*; the 3 temporal ones answer *how it evolves* — together they separate the 5 drift types with no labels.

Rule-based Logic to Identify Drift Types



Backup: IDW-MMD vs Optimally-Weighted MMD (Bharti et al.)

Question: “How is IDW-MMD different from the Optimally-Weighted MMD of Bharti et al. (2023)?”

Bharti et al. (OW-MMD)

- Domain: **likelihood-free inference** (ABC).
- Weights chosen by optimising against a target posterior.
- Requires solving a constrained convex optimisation per call.
- Not designed for streaming drift detection.

IDW-MMD (this work)

- Domain: **drift detection** in fast data streams.
- Weights from a closed-form heuristic: $w_i \propto 1/(\sqrt{d_i} + \epsilon)$.
- Cost: $O(n^2)$ for the kernel matrix, $O(n)$ extra for the weights — no optimisation step.
- Same *weighted MMD* family, different objective and very different runtime.

Summary: both are weighted MMD variants, but they target different problems (ABC vs streaming drift), use different weight functions, and have very different runtime profiles.



Backup: “F1 close to MMD baseline — where is the contribution?”

Question: “SE-CDT’s F1 (0.531) is only marginally above MMD (0.524). Where is the gain?”

F1 alone hides the trade-off:

	SE-CDT	MMD	KS
F1	0.531	0.524	0.335
FP	9.6	10.6	27.2
Recall	0.737	0.732	0.775
Precision	0.432	0.426	0.224

In a production setting:

- Frequent FPs → **alert fatigue**, lost trust in the monitor.
- Investigating one false alarm typically costs more than missing one event.
- A miscalibrated permutation-style test inflates either FPs or misses.

What SE-CDT actually adds:

- 1 **Single-test calibrated near $\alpha = 0.05$** — Standard MMD / IDW-MMD Type-I error 0.025–0.075; the SE-CDT composite is anti-conservative (0.080–0.140), reported honestly as a pipeline limitation.
- 2 **Classification module on top of detection** — gives a drift type, not just a flag.
- 3 **Modest runtime speedup over ShapeDD** thanks to the Gamma-null p-value (no kernel re-build per shuffle).
- 4 Stays **fully unsupervised** at runtime.



Backup: Why does Concept Memory matter for Recurrent drift?

Question: “Recurrent drift requires concept matching — how does Concept Memory provide it?”

The detection challenge:

- A Recurrent stream is a sequence of *repeated* concept switches; each switch produces an MMD peak that, in isolation, looks like a Sudden drift.
- Without memory of past concepts, repeated visits are scored independently $\rightarrow$ Recurrent collapses into the Sudden branch.

How Concept Memory resolves it:

- 1 During every *stable period*, take a snapshot of the current data window and store it in a ring buffer of size $K = 8$.
- 2 When a new drift fires, compute Standard MMD between the new post-drift window and every snapshot.
- 3 If the smallest $\text{MMD} < \mu = 0.15 \rightarrow$ override the

Why $\mu = 0.15$?

$\mu = 0.05$ falls inside the MMD noise floor (false matches between independent stationary windows). An empirical sweep on H0 calibration data shows $\mu = 0.15$ separates “revisited concept” from “new concept” reliably.

Remaining limitations

Two failure modes remain on the harder Recurrent variants:

- Recurrent events that are *too far apart* push the original snapshot out of the ring buffer.
- Repeated Gradual streams saturate after a few events, so the post-drift window matches

Backup: “CDT-MSW has higher CAT and SUB — is SE-CDT worse?”

Question: “CDT-MSW reports CAT 87.0% and SUB 86.7%, both higher than SE-CDT. Why frame SE-CDT as competitive?”

Side-by-side numbers

	SE-CDT	CDT-MSW
Labels	None	Required
CAT	80.1%	87.0%
SUB	55.3%	86.7%
Recall	92.6%	82.5%
Effective	74.2%	71.8%

Effective = CAT $\times$ Recall (probability of seeing a drift *and* typing it correctly). CDT-MSW Recall = $1 - \overline{\text{MDR}}$ from paper Table 5.

Reading the numbers in context:

- CDT-MSW reports CAT and SUB *conditionally* on drifts it detects; with paper recall $\approx 82.5\%$ it still misses $\sim 17\%$ of events.
- SE-CDT detects 92.6% of drifts and types every one of them without any labels.
- End-to-end, SE-CDT lands 2.4pp above CDT-MSW on effective accuracy (74.2% vs 71.8%); the contribution is reaching this on **unsupervised** $P(X)$.

Trade-off

CDT-MSW is more accurate per detected event because it has access to ground truth labels. SE-CDT

Backup: “Why not deep learning?”

Question: “Couldn’t a deep learning model solve this more accurately?”

Deep models for drift detection — limitations

- Need a large labelled dataset to train — the very thing we don’t have on a streaming pipeline.
- Higher inference latency per window than a kernel computation.
- No closed-form theoretical bound on the test statistic.
- Decision is hard to explain to a domain expert.

Why MMD / ShapeDD fits this problem

- Fully unsupervised — no training run before deployment.
- Theoretical guarantees: Triangle Shape Theorem, Type-I error from H_0 calibration.
- Decision rule is human-readable (a small tree of geometric features).
- One mechanism works across $P(X)$ shifts of arbitrary type.

Future direction: keep the kernel framework but plug in a learned feature extractor for high-dimensional inputs (Chapter 6).



Backup: behaviour on mild drift


Question: “On the mild-drift scenario (Random Mild, $\delta = 0.125$) SE-CDT scored $F1 = 0.511$ vs MMD $= 0.563$ — where does the gap come from?”

Cause:

- IDW down-weights points in dense (stable) regions.
- Mild drift lives mostly inside those regions, shortening the IDW peak.
- Net effect: some peaks fall under threshold, reducing recall.

This is the intended trade-off:

- IDW-MMD biases toward fewer false alarms (FP ~ 9.6 /run).
- Cost: ~ 5 pp F1 drop on the smallest shifts.



Possible mitigations (future work)



Backup: Permutation test vs Gamma-null p-value

Question: “Removing the permutation test — is the Gamma-null p-value still trustworthy?”

Permutation test (original ShapeDD)

- Reshuffles indices many times per window.
- Empirical null, no distributional assumption.
- Each shuffle re-builds the kernel matrix $\rightarrow$ heavy compute.

Gamma-null p-value (IDW-MMD)


- Moment-matched Gamma fit to the null MMD distribution (Gretton et al. 2009).
- Estimates the first two moments from a small number of index shuffles, reusing the kernel matrix.
- Complexity $O(n^2)$ for the kernel plus a few cheap shuffles.

Verification: The detection F_1 remains comparable to the permutation baseline; the H_0 -calibration on three stationary distributions reports Standard MMD and IDW-MMD inside the 95% CI of $\alpha = 0.05$ (0.025–0.075), while the SE-CDT composite is anti-conservative at 0.080–0.140 (up to $\approx 2.8 \times \alpha$), reflecting peak-selection and sliding-window dependence rather than the per-candidate test itself. End-to-end runtime is slightly faster than the ShapeDD baseline.



Backup: “Subcategory accuracy 55.3% — is it close to random?”

Question: “With 5 classes, is $\sim 60\%$ close to random guessing?”



It is well above random:

- Random baseline across 5 classes is 20%.
- Micro SUB = 55.3% is $\sim 2.8\times$ random and well above the 95% CI of a random predictor ($\sim 23\%$).
- Performance is **very uneven** across classes:

Drift Type	Accuracy
------------	----------

Sudden	82.4%
--------	-------


Incremental	30.3%
-------------	-------

Recurrent	71.5%
-----------	-------



Backup: Computational Overhead & Runtime Analysis

Question: “What is the computational overhead of IDW-MMD compared to standard methods?”




Complexity & Efficiency:

- **Standard MMD / IDW-MMD:** $O(n^2)$ matrix computation, but IDW computes weights once per window step.
- **Gamma-null Approximation** bypasses the permutation test entirely, running in constant extra shuffles.
- **ShapeDD-IDW** is roughly **7 $\times$** faster than the original permutation-based ShapeDD (0.59s vs 4.26s per 10k samples).

Streaming Suitability

In online deployments, the detector easily processes streams at over **16k samples/sec**, making its latency negligible in streaming Kafka pipelines.

Backup: Kafka Deployment — System Resource Usage



Memory and CPU stay stable throughout the run — SE-CDT's streaming overhead is negligible.

Backup: Why are Gradual (27.2%) and Incremental (30.3%) low?

Question: “The two slow types are the weak point — what exactly goes wrong?”

Two systematic confusions:

- 1 **Gradual** $\leftrightarrow$ **Incremental overlap:** Both land in ambiguous Variance Ratio bands ($\sim 47\%$ of Incremental read as Gradual).
- 2 **Saturation** $\rightarrow$ **Recurrent:** In repeated slow drifts, $P(X)$ stops changing. The post-drift window matches stored snapshots, labelling it Recurrent ($\sim 57\%$ of Gradual events).

Why CAT is unaffected

Both confusions stay *inside* PCD, preserving TCD/PCD accuracy (CAT = 80.1%). Confusing Gradual with Incremental is low-cost for adaptation.

Fix (future work)

Replace fixed VR thresholds with learned classifiers and separate Concept-Memory match step from shape classification.

Backup: “Is the adaptation gain Sudden-biased?”

Question: “Type-Specific wins +17pp — but the classifier is only $\sim 55\%$ on subtypes. Does the gain depend on the easy scenarios?”

Honest answer — yes, partly:

- The three adaptation scenarios (Stepping, Sudden, Mixed) are all **Sudden-dominated**, exactly where the classifier is strongest (Sudden 82.4%).
- So routing is correct *most of the time in this region*, and any reset on a true sudden drift helps regardless.
- We have **not** tested adaptation on Gradual/Incremental-heavy streams, where subtype accuracy drops to 27–30%.

Why the gain still holds

The category decision (TCD vs PCD, 80.1%) drives the adaptation policy, not the exact subtype. TCD→reset and PCD→fine-tune is the right call even when the subtype within PCD is wrong — the operational cost of confusing Gradual with Incremental is small (same update family).

Scope claim

We claim the adaptation benefit *for the sudden/recurrent regime*, not universally. Gradual/Incremental-heavy adaptation is listed as future work (Ch. 6).

